

In [1]: *# Import Libraries*

```
import pandas as pd
import numpy as np
import tensorflow as tf
import h5py
from tqdm import tqdm
import matplotlib.pyplot as plt
import seaborn as sns
```

2023-12-04 23:16:19.861025: I tensorflow/core/platform/cpu\_feature\_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.  
To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

In [2]: *# Define models*

```
clean_test_data = r'test.h5'
poisoned_test_data = r'bd_test.h5'
clean_val_data = r'valid.h5'
poisoned_val_data = r'bd_valid.h5'
badnet_model = r'bd_net.h5'
```

In [3]: *# Define GoodNet Class*

```
class G(tf.keras.Model):
    def __init__(self, B_path, B_prime_path):
        super(G, self).__init__()
        self.B = tf.keras.models.load_model(B_path)
        self.B_prime = tf.keras.models.load_model(B_prime_path)
        self.backdoor_class_index = self.B.layers[-1].output_shape[-1]

    def predict(self, data):
        y = np.argmax(self.B.predict(data), axis=1)
        y_prime = np.argmax(self.B_prime.predict(data), axis=1)
        pred = np.zeros(data.shape[0])
        for i in range(data.shape[0]):
            if y[i] == y_prime[i]:
                pred[i] = y[i]
            else:
                pred[i] = self.backdoor_class_index
        return pred
```

In [4]: *# Model Evaluation Function*

```
def data_loader(filepath):
    data = h5py.File(filepath, 'r')
    x_data = np.array(data['data'])
    y_data = np.array(data['label'])
    x_data = x_data.transpose((0,2,3,1))
    return x_data, y_data

def evaluate_model(clean_data_filename, poisoned_data_filename, model_input):
    cl_x_test, cl_y_test = data_loader(clean_data_filename)
    bd_x_test, bd_y_test = data_loader(poisoned_data_filename)
```

```

# Determine the type of model_input and load or use the model accordingly
if isinstance(model_input, str):
    model = tf.keras.models.load_model(model_input)
    predict_function = lambda x: np.argmax(model.predict(x), axis=1)
elif isinstance(model_input, G):
    model = model_input
    predict_function = lambda x: model.predict(x)
else: # Assuming model_input is a Keras model
    model = model_input
    predict_function = lambda x: np.argmax(model.predict(x), axis=1)

# Evaluate the model
cl_label_prd = predict_function(cl_x_test)
cl_clf_acc = np.mean(np.equal(cl_label_prd, cl_y_test)) * 100
print('Classification Accuracy on Clean Data:', cl_clf_acc)

bd_label_prd = predict_function(bd_x_test)
at_suc_rt = np.mean(np.equal(bd_label_prd, bd_y_test)) * 100
print('Attack Success Rate using Poisoned Data:', at_suc_rt)

return cl_clf_acc, at_suc_rt

```

```

In [5]: # Get baseline stats for model

bl_cl_clf_acc, bl_at_suc_rt = evaluate_model(clean_val_data, poisoned_val_data)

361/361 [=====] - 2s 5ms/step
Classification Accuracy on Clean Data: 98.64899974019225
361/361 [=====] - 2s 6ms/step
Attack Success Rate using Poisoned Data: 100.0

```

```

In [6]: # Visualize the Model Structure

model = tf.keras.models.load_model(badnet_model)
print(model.summary())

```

Model: "model\_1"

Layer (type) to	Output Shape	Param #	Connected
input (InputLayer)	[(None, 55, 47, 3)]	0	[]
conv_1 (Conv2D)	(None, 52, 44, 20)	980	['input[0][0]']
pool_1 (MaxPooling2D)	(None, 26, 22, 20)	0	['conv_1[0][0]']
conv_2 (Conv2D)	(None, 24, 20, 40)	7240	['pool_1[0][0]']
pool_2 (MaxPooling2D)	(None, 12, 10, 40)	0	['conv_2[0][0]']
conv_3 (Conv2D)	(None, 10, 8, 60)	21660	['pool_2[0][0]']
pool_3 (MaxPooling2D)	(None, 5, 4, 60)	0	['conv_3[0][0]']
conv_4 (Conv2D)	(None, 4, 3, 80)	19280	['pool_3[0][0]']
flatten_1 (Flatten)	(None, 1200)	0	['conv_4[0][0]']
flatten_2 (Flatten)	(None, 960)	0	['flatten_1[0][0]']
fc_1 (Dense)	(None, 160)	192160	['flatten_2[0][0]']
fc_2 (Dense)	(None, 160)	153760	['fc_1[0][0]']
add_1 (Add)	(None, 160)	0	['fc_1[0][0]', 'fc_2[0][0]']
activation_1 (Activation)	(None, 160)	0	['add_1[0][0]']
output (Dense)	(None, 1283)	206563	['activation_1[0][0]']
Total params: 601643 (2.30 MB)			
Trainable params: 601643 (2.30 MB)			
Non-trainable params: 0 (0.00 Byte)			

None

```
In [7]: # Clear the current Tensorflow Keras session
tf.keras.backend.clear_session()
```

```
In [8]: # Visualize the clean data

x_data, y_data = data_loader(clean_val_data)
figure = plt.figure(figsize=(10,8))
cols, rows = 3,3

for i in range(1, cols*rows+1):
    index = np.random.randint(x_data.shape[0], size=1)
    img, label = (x_data[index], y_data[index])
    figure.add_subplot(rows, cols, i)
    plt.title("True Label: {}".format(label))
    plt.axis("off")
    plt.imshow(img[0]/255)
plt.show()
```


True Label: [812.]



True Label: [1186.]



True Label: [218.]



True Label: [66.]



True Label: [1109.]



True Label: [282.]



True Label: [409.]



True Label: [273.]



True Label: [66.]



```
In [9]: # Visualize the poisoned data
```

```

x_data, y_data = data_loader(poisoned_val_data)
figure = plt.figure(figsize=(10,8))
cols, rows = 3,3

for i in range(1, cols*rows+1):
    index = np.random.randint(x_data.shape[0], size=1)
    img, label = (x_data[index], y_data[index])
    figure.add_subplot(rows, cols, i)
    plt.title("True Label: {}".format(label))
    plt.axis("off")
    plt.imshow(img[0]/255)
plt.show()

```



```

In [10]: # Create a copy of the BadNet model

model = tf.keras.models.load_model(badnet_model)
source_optimizer = model.optimizer
source_loss = model.loss
source_metrics = model.metrics
bd_model_cpy = tf.keras.models.clone_model(model)
bd_model_cpy.set_weights(model.get_weights())

```

```

In [11]: # Pruning models as a defence against poisoned model

```

```

cl_clf_acc_arr = []
at_suc_rt_arr = []
saved_model = np.zeros(3,dtype=bool)

output_layer = bd_model_cpy.get_layer('pool_3').output
seq = np.argsort(np.mean(tf.keras.models.Model(inputs = bd_model_cpy.input, ou
weight_0 = bd_model_cpy.layers[5].get_weights()[0]
bias_0 = bd_model_cpy.layers[5].get_weights()[1]

for channel_index in tqdm(seq):
    weight_0[:, :, :, channel_index] = 0
    bias_0[channel_index] = 0
    bd_model_cpy.layers[5].set_weights([weight_0, bias_0])
    cl_clf_acc, at_suc_rt = evaluate_model(clean_val_data, poisoned_val_data, l
    if (bl_cl_clf_acc - cl_clf_acc >= 2 and not saved_model[0]):
        print("Model Saved at 2% drop in classification accuracy.")
        bd_model_cpy.compile(optimizer=source_optimizer,
                               loss=source_loss,
                               metrics=source_metrics)
        bd_model_cpy.save('model_X_2.h5')
        saved_model[0] = 1
    if (bl_cl_clf_acc - cl_clf_acc >= 4 and not saved_model[1]):
        print("Model Saved at 4% drop in classification accuracy.")
        bd_model_cpy.compile(optimizer=source_optimizer,
                               loss=source_loss,
                               metrics=source_metrics)
        bd_model_cpy.save('model_X_4.h5')
        saved_model[1] = 1
    if (bl_cl_clf_acc - cl_clf_acc >= 10 and not saved_model[2]):
        print("Model Saved at 10% drop in classification accuracy.")
        bd_model_cpy.compile(optimizer=source_optimizer,
                               loss=source_loss,
                               metrics=source_metrics)
        bd_model_cpy.save('model_X_10.h5')
        saved_model[2] = 1
    cl_clf_acc_arr.append(cl_clf_acc)
    at_suc_rt_arr.append(at_suc_rt)
    tf.keras.backend.clear_session()

```

361/361 [=====] - 2s 5ms/step

0%|

| 0/60 [00:00<?, ?it/s]

361/361 [=====] - 3s 7ms/step

Classification Accuracy on Clean Data: 98.64899974019225

361/361 [=====] - 2s 6ms/step

Attack Success Rate using Poisoned Data: 100.0

2%|■

| 1/60 [00:08<08:07, 8.26s/it]

361/361 [=====] - 2s 6ms/step

Classification Accuracy on Clean Data: 98.64899974019225

361/361 [=====] - 2s 6ms/step

Attack Success Rate using Poisoned Data: 100.0

3%|■■

| 2/60 [00:16<07:42, 7.98s/it]

361/361 [=====] - 2s 6ms/step

Classification Accuracy on Clean Data: 98.64899974019225

361/361 [=====] - 2s 6ms/step

Attack Success Rate using Poisoned Data: 100.0

5%|██████████  
| 3/60 [00:23<07:31, 7.92s/it]  
361/361 [=====] - 2s 6ms/step  
Classification Accuracy on Clean Data: 98.64899974019225  
361/361 [=====] - 2s 6ms/step  
Attack Success Rate using Poisoned Data: 100.0

7%|██████████  
| 4/60 [00:31<07:20, 7.86s/it]  
361/361 [=====] - 2s 6ms/step  
Classification Accuracy on Clean Data: 98.64899974019225  
361/361 [=====] - 2s 6ms/step  
Attack Success Rate using Poisoned Data: 100.0

8%|██████████  
| 5/60 [00:39<07:13, 7.88s/it]  
361/361 [=====] - 2s 6ms/step  
Classification Accuracy on Clean Data: 98.64899974019225  
361/361 [=====] - 2s 6ms/step  
Attack Success Rate using Poisoned Data: 100.0

10%|██████████  
| 6/60 [00:47<07:03, 7.85s/it]  
361/361 [=====] - 2s 6ms/step  
Classification Accuracy on Clean Data: 98.64899974019225  
361/361 [=====] - 2s 6ms/step  
Attack Success Rate using Poisoned Data: 100.0

12%|██████████  
| 7/60 [00:55<07:04, 8.00s/it]  
361/361 [=====] - 3s 7ms/step  
Classification Accuracy on Clean Data: 98.64899974019225  
361/361 [=====] - 2s 6ms/step  
Attack Success Rate using Poisoned Data: 100.0

13%|██████████  
| 8/60 [01:04<07:04, 8.17s/it]  
361/361 [=====] - 2s 6ms/step  
Classification Accuracy on Clean Data: 98.64899974019225  
361/361 [=====] - 2s 6ms/step  
Attack Success Rate using Poisoned Data: 100.0

15%|██████████  
| 9/60 [01:12<06:51, 8.06s/it]  
361/361 [=====] - 2s 6ms/step  
Classification Accuracy on Clean Data: 98.64899974019225  
361/361 [=====] - 2s 6ms/step  
Attack Success Rate using Poisoned Data: 100.0

17%|██████████  
| 10/60 [01:19<06:37, 7.96s/it]  
361/361 [=====] - 2s 6ms/step  
Classification Accuracy on Clean Data: 98.64899974019225  
361/361 [=====] - 2s 6ms/step  
Attack Success Rate using Poisoned Data: 100.0

18%|██████████  
| 11/60 [01:27<06:30, 7.97s/it]  
361/361 [=====] - 2s 6ms/step  
Classification Accuracy on Clean Data: 98.64899974019225  
361/361 [=====] - 2s 6ms/step  
Attack Success Rate using Poisoned Data: 100.0

20%|██████████  
| 12/60 [01:35<06:19, 7.90s/it]





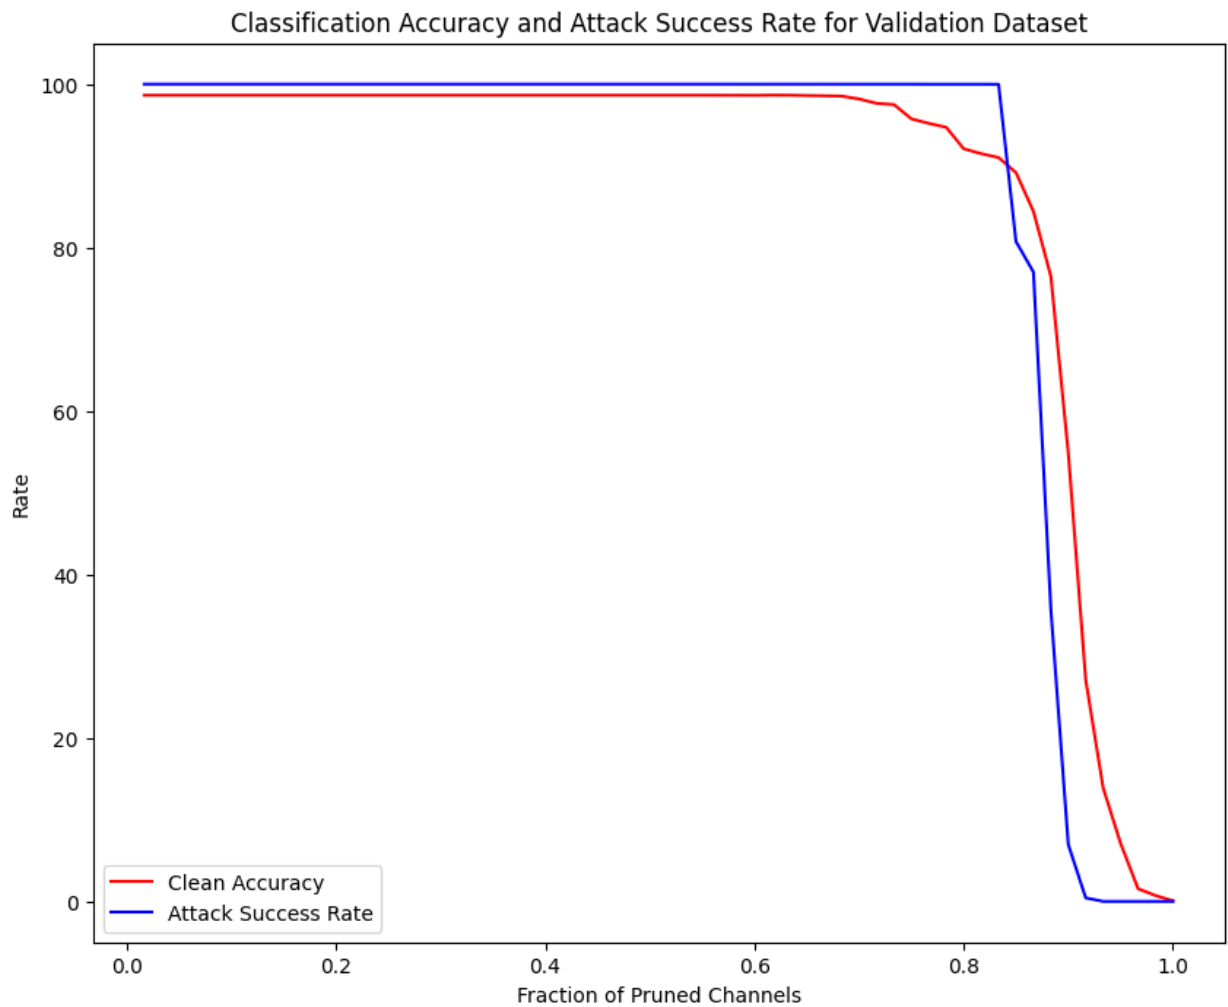












In [13]: *# Clear the current Tensorflow Keras session*

```
tf.keras.backend.clear_session()
```

In [14]: *# Define Pruned models*

```
prune_model_X_2 = 'model_X_2.h5'
prune_model_X_4 = 'model_X_4.h5'
prune_model_X_10 = 'model_X_10.h5'
```

In [15]: *# Define GoodNet(Repaired) Models*

```
G_model_X_2 = G(badnet_model, prune_model_X_2)
G_model_X_4 = G(badnet_model, prune_model_X_4)
G_model_X_10 = G(badnet_model, prune_model_X_10)
```

In [16]: *# Evaluate GoodNet(Repaired) Models*

```
cl_clf_acc_X_2, at_suc_rt_X_2 = evaluate_model(clean_test_data, poisoned_test_data, G_model_X_2)
cl_clf_acc_X_4, at_suc_rt_X_4 = evaluate_model(clean_test_data, poisoned_test_data, G_model_X_4)
cl_clf_acc_X_10, at_suc_rt_X_10 = evaluate_model(clean_test_data, poisoned_test_data, G_model_X_10)
```

```

401/401 [=====] - 2s 5ms/step
401/401 [=====] - 3s 6ms/step
Classification Accuracy on Clean Data: 95.74434918160561
401/401 [=====] - 3s 6ms/step
401/401 [=====] - 3s 6ms/step
Attack Success Rate using Poisoned Data: 100.0
401/401 [=====] - 3s 7ms/step
401/401 [=====] - 2s 6ms/step
Classification Accuracy on Clean Data: 92.1278254091972
401/401 [=====] - 3s 7ms/step
401/401 [=====] - 3s 7ms/step
Attack Success Rate using Poisoned Data: 99.98441153546376
401/401 [=====] - 3s 6ms/step
401/401 [=====] - 3s 7ms/step
Classification Accuracy on Clean Data: 84.3335931410756
401/401 [=====] - 3s 6ms/step
401/401 [=====] - 3s 7ms/step
Attack Success Rate using Poisoned Data: 77.20966484801247

```

```

In [17]: # Table of GoodNet(Repaired) model evaluation

cl_clf_acc_tb = [cl_clf_acc_X_2, cl_clf_acc_X_4, cl_clf_acc_X_10]
at_suc_rt_tb = [at_suc_rt_X_2, at_suc_rt_X_4, at_suc_rt_X_10]
data = {
    "Classification Accuracy on Clean Data": cl_clf_acc_tb,
    "Attack Success Rate using Poisoned Data": at_suc_rt_tb,
    "Model": ["GoodNet_2%", "GoodNet_4%", "GoodNet_10%"]
}
df = pd.DataFrame(data)
df.set_index('Model')

```

```

Out[17]:

```

	Classification Accuracy on Clean Data	Attack Success Rate using Poisoned Data
Model		
GoodNet_2%	95.744349	100.000000
GoodNet_4%	92.127825	99.984412
GoodNet_10%	84.333593	77.209665

```

In [18]: # Graph of GoodNet(Repaired) model evaluation

opacity = 1
bar_width = 0.25

figure = plt.figure(figsize=(10,8))
plt.xlabel('Drop in Classification Accuracy of Clean Data')
plt.ylabel('Percentage')

plt.xticks(range(len(cl_clf_acc_tb)),('2', '4', '10'))
bar1 = plt.bar(np.arange(len(cl_clf_acc_tb)) + bar_width, cl_clf_acc_tb, bar_w
bar2 = plt.bar(range(len(at_suc_rt_tb)),at_suc_rt_tb, bar_width, align='center

for rect in bar1 + bar2:
    height = rect.get_height()
    plt.text(rect.get_x() + rect.get_width() / 2.0, height, f'{height:.02f}', l

plt.legend(bbox_to_anchor=(1.4, 1))

```

```
plt.tight_layout()
plt.title('Performance of the GoodNet(Repaired) Model')
sns.despine()
plt.show()
```

